

Introduction to R

Using basic functions in R

Inspecting function documentation

There are many functions in base R that you can use (in a later tutorial, we will discuss how to create your own functions!). In order to view documentation for a function in R, you can type `?function_name` or search for the function name in the help tab. As practice, let's explore the `matrix()` and `mean()` functions in R by typing `?matrix` and `?mean` into the console. Uncomment (by deleting the `#s`) the code below and run it in the console to inspect the documentation.

Under usage for the `matrix()` function, you should see the following:

```
matrix(data = NA, nrow = 1, ncol = 1, byrow = FALSE, dimnames = NULL)
```

The `matrix` function has 5 distinct parameters. All of them have default values. For example, if you don't put any data into the function, the resulting matrix will be made up of NA values (this is how missing data is generally coded in R). Additionally, the function will have 1 row and 1 column by default and data will be filled by columns because `byrow=FALSE` by default. The matrix will not have any row or column names because `dimnames=NULL` by default.

```
#?matrix  
#?mean
```

Because all parameters of the `matrix()` function have defaults, we could call `matrix()` with no inputs and we would get a 1x1 matrix with NA values and no dimension names. We can also pick and choose whatever parameters we do want to fill in and ignore anything that we want to leave as defaults.

```
# Call matrix() with no inputs  
matrix()
```

```
##      [,1]  
## [1,]  NA
```

```
# Make a 2x3 matrix of NAs  
matrix(nrow=2, ncol=3)
```

```
##      [,1] [,2] [,3]  
## [1,]  NA  NA  NA  
## [2,]  NA  NA  NA
```

In contrast, the `mean()` function has some required parameters. When you type `?mean` into the console, you should see the following:

```
mean(x, trim = 0, na.rm = FALSE, ...)
```

The parameter `x` has no default value; you must specify the values that you want to take the mean of. The other parameters are given defaults: `trim=0` and `na.rm=FALSE`. If you read the descriptions for these parameters, you will see that `trim` allows you to calculate a trimmed mean (i.e., eliminate some proportion of extreme values before calculating the mean) and `na.rm` allows you to remove missing (NA) values before calculating the mean. By default, R will not do any trimming and NAs will not be removed. This can cause issues (see below):

```
# Create some data and save it as data1 and data2  
data1 <- c(1,2,3,4,7,NA)  
data2 <- c(1,2,3,4,7)
```

```
# Calculate the mean of data1 and data2  
# Note that data1 has a mean of NA because there was an NA value that was not removed  
mean(data1)
```

```
## [1] NA
mean(data2)

## [1] 3.4
# Now explicitly set na.rm=TRUE and recalculate the mean of data1. Now we get the same as data2
mean(data1, na.rm=TRUE)

## [1] 3.4
```

Calling functions in R

When using functions in R, you can make it clear which inputs refer to which parameters by either a) using the order of the parameters specified in the usage or b) naming them explicitly.

```
# Note that these two calls to the mean() function will return the same output
mean(data1, .2, TRUE)

## [1] 3
mean(x=data1, trim=.2, na.rm=TRUE)

## [1] 3
# However, the following code will give an error because "TRUE" is not a valid input to trim,
# Which is the second parameter listed in the usage
# Mean(data1,TRUE)

# To specify na.rm=TRUE but leave the trim=0 default as is, we simply do the following:
mean(data1, na.rm=TRUE)

## [1] 3.4
# (Note that, since data1 still matches to x, which is the first parameter in usage,
# we do not need to specify x=data1)
```

Data types and structures

This section covers four basic data types in R: characters, numerics, integers, and logicals. This section covers four basic ways to store data in R: vectors, data frames, matrices, and lists.

You can learn about these data types below

Numerics and integers

```
# To store data in some variable name, use either = or <-
# To save the number 5 as "number":
number <- 5
number = 5 #does the same thing

# Print number:
print(number)

## [1] 5
# Find out the class of number:
class(number)
```

```
## [1] "numeric"
# Change number to an integer and re-save it as number2:
number2 <- as.integer(number)

# Inspect the class of number2 to see that it is an integer:
class(number2)

## [1] "integer"
```

Characters

```
# Save a character string in a variable named message
message <- "welcome"

# Print message:
print(message)

## [1] "welcome"

# Inspect the class of message:
class(message)

## [1] "character"
```

Logicals

Logical data is either TRUE or FALSE

In R, TRUE=1 and FALSE=0

```
# Save the logical TRUE as a variable called outcome:
outcome <- TRUE

# Print outcome
print(outcome)

## [1] TRUE

# Inspect class of outcome
class(outcome)

## [1] "logical"

# Note that, weirdly, outcome+outcome=2
outcome+outcome

## [1] 2
```

In R, you can test a statement to see if it is TRUE or FALSE. Note that R allows you to make comparisons across variable types: integers may be compared to numerics and logicals may be compared to integers/numerics. For characters, comparatives are assessed using alphabetical order (letters earlier in the alphabet are “smaller”):

- 1) == means “is equal to”
- 2) != means “is not equal to”
- 3) > means “greater than”; >= means “greater than or equal to”
- 4) < means “less than; <= means “less than or equal to”

```
# Is 5 equal to 3?
5==3
```

```
## [1] FALSE
# Is 5 not equal to 3?
5!=3

## [1] TRUE
# Is 5 less than 3?
5<3

## [1] FALSE
# Is 5 greater than 3?
5>3

## [1] TRUE
# Is 5 greater than 5?
5>5

## [1] FALSE
# Is 5 greater than or equal to 5?
5>=5

## [1] TRUE
# Is 5 equal to 5?
5==5

## [1] TRUE
# Is "hello" equal to "hello"?
"hello" == "hello"

## [1] TRUE
# Is "hello" equal to "goodbye"?
"hello" == "goodbye"

## [1] FALSE
# Is "hello" greater than "goodbye"? (in other words is "hello" after "goodbye" alphabetically?)
"hello">"goodbye"

## [1] TRUE
# Is TRUE == 1?
TRUE==1

## [1] TRUE
# Is FALSE==0?
FALSE==0

## [1] TRUE
```

Vectors

```
# The easiest way to create a vector is by using the c() function
vec1 <- c(2,3,4,5)
print(vec1)

## [1] 2 3 4 5
```

```
# Note that, if you include multiple data types in a vector, R will change all values to the same type
vec2 <- c(7,3)
vec2
```

```
## [1] 7 3
```

```
class(vec2)
```

```
## [1] "numeric"
```

```
vec3 <- c(7,3,"hello")
vec3
```

```
## [1] "7"      "3"      "hello"
```

```
class(vec3)
```

```
## [1] "character"
```

```
# For values in a row, we can also use a colon:
```

```
vec3 <- 2:5
print(vec3)
```

```
## [1] 2 3 4 5
```

```
# We can use the c() function to combine pre-saved vectors:
```

```
vec4 <- c(vec1,vec3)
print(vec4)
```

```
## [1] 2 3 4 5 2 3 4 5
```

```
# Use the length function to find the length of a vector
```

```
length(vec4)
```

```
## [1] 8
```

Here are some other useful shortcuts for creating vectors in R:

```
# Use the rep() function to repeat values. Inspect the following to see how it works!
```

```
rep(x=0,times=5) #create a vector of 5 zeros
```

```
## [1] 0 0 0 0 0
```

```
rep(x=c(1,2,3),times=5) #create a vector with five repeats of 1,2,3
```

```
## [1] 1 2 3 1 2 3 1 2 3 1 2 3 1 2 3
```

```
rep(c(1,2,3),each=2,times=5) #repeat each value in 1,2,3 twice, then repeat that 5 times
```

```
## [1] 1 1 2 2 3 3 1 1 2 2 3 3 1 1 2 2 3 3 1 1 2 2 3 3
```

```
# Use the seq function to create a sequence of values
```

```
seq(from=1, to=5, by=.5) #create a vector with values from 1-5, incrementing by .5
```

```
## [1] 1.0 1.5 2.0 2.5 3.0 3.5 4.0 4.5 5.0
```

```
seq(from=1, to=5, length.out=17) #create a vector with 17 equally spaced values going from 1 to 5
```

```
## [1] 1.00 1.25 1.50 1.75 2.00 2.25 2.50 2.75 3.00 3.25 3.50 3.75 4.00 4.25
```

```
## [15] 4.50 4.75 5.00
```

Matrices

As shown above, an m by n matrix can be created in R using the `matrix()` function
Here are some examples

```
A <- matrix(2, nrow=3, ncol=3)
print(A)
```

```
##      [,1] [,2] [,3]
## [1,]    2    2    2
## [2,]    2    2    2
## [3,]    2    2    2
```

```
B <- matrix(c(1,2,5,3,4,0,2,1,5), nrow=3, ncol=3, byrow=TRUE)
print(B)
```

```
##      [,1] [,2] [,3]
## [1,]    1    2    5
## [2,]    3    4    0
## [3,]    2    1    5
```

```
# Print the dimensions of a matrix (number of rows followed by number of columns)
dim(A)
```

```
## [1] 3 3
```

```
# Matrix multiplication
A %*% B
```

```
##      [,1] [,2] [,3]
## [1,]   12   14   20
## [2,]   12   14   20
## [3,]   12   14   20
```

```
# Element-wise multiplication
A * B
```

```
##      [,1] [,2] [,3]
## [1,]    2    4   10
## [2,]    6    8    0
## [3,]    4    2   10
```

```
# Element-wise addition
A+B
```

```
##      [,1] [,2] [,3]
## [1,]    3    4    7
## [2,]    5    6    2
## [3,]    4    3    7
```

```
# Transpose of a matrix
t(A)
```

```
##      [,1] [,2] [,3]
## [1,]    2    2    2
## [2,]    2    2    2
## [3,]    2    2    2
```

```
# Inverse of a matrix
solve(B)
```

```
##      [,1]      [,2]      [,3]
## [1,] -0.5714286  0.14285714  0.57142857
```

```
## [2,] 0.4285714 0.14285714 -0.42857143
## [3,] 0.1428571 -0.08571429 0.05714286
```

Data Frames

Data frames are generally used to store tabular data. Data frames are composed of same-length vectors; these vectors can be of differing data types. In general, when you read a .csv data file into R, it will be saved as a data frame.

We can create a data frame in R as follows:

```
# Create a fake dataset called example_data
example_data <- data.frame(ID_Num = c(1:10),
                           Age = rep(24:28, each=2),
                           State = c(rep("New Jersey", 5), rep("New York", 5)))

# Change row names of the data frame (some made up names)
rownames(example_data) <- c("Sarah", "Mike", "Drew", "Eric", "Maria",
                             "Lindsey", "Mark", "Jenny", "Sophie", "Paul")

# Print the data frame
example_data
```

```
##      ID_Num Age      State
## Sarah      1  24 New Jersey
## Mike       2  24 New Jersey
## Drew       3  25 New Jersey
## Eric       4  25 New Jersey
## Maria      5  26 New Jersey
## Lindsey    6  26  New York
## Mark       7  27  New York
## Jenny      8  27  New York
## Sophie     9  28  New York
## Paul     10  28  New York
```

The following R code outlines a few ways to inspect data in a data frame.

```
# Get dimensions (same as matrices)
dim(example_data)
```

```
## [1] 10 3
```

```
# Get number of columns
ncol(example_data)
```

```
## [1] 3
```

```
# Get number of rows
nrow(example_data)
```

```
## [1] 10
```

```
# Get summaries of the columns
summary(example_data)
```

```
##      ID_Num      Age      State
## Min.   : 1.00  Min.   :24  New Jersey:5
## 1st Qu.: 3.25  1st Qu.:25  New York :5
## Median : 5.50  Median :26
```

```
## Mean : 5.50 Mean :26
## 3rd Qu.: 7.75 3rd Qu.:27
## Max. :10.00 Max. :28

# Access a single column of the data frame using $
example_data$Age

## [1] 24 24 25 25 26 26 27 27 28 28

# Inspect row names
rownames(example_data)

## [1] "Sarah" "Mike" "Drew" "Eric" "Maria" "Lindsey" "Mark"
## [8] "Jenny" "Sophie" "Paul"

# Inspect column names
colnames(example_data)

## [1] "ID_Num" "Age" "State"
```

Lists

Lists enable multiple data types or data sets to be stored in a simple object. For example, a list could have a data frame as its first element, a vector as its second element, and a character string as its third element.

```
# Save vector vec1, matrix A, and vector vec2 in a list called example list
example_list <- list(vec1, A, vec2)

# Print example_list
example_list

## [[1]]
## [1] 2 3 4 5
##
## [[2]]
##      [,1] [,2] [,3]
## [1,]    2    2    2
## [2,]    2    2    2
## [3,]    2    2    2
##
## [[3]]
## [1] 7 3
```

Indexing

In R, indices start at 1, not 0 as in other languages. For example, the index of the 3rd element in a vector is 3.

Using indices to extract elements in a vector

We can use indices enclosed in square brackets in order to extract data from a vector as follows:

```
# This R chunk uses vector vec4 from above
# Re-print vec4
vec4

## [1] 2 3 4 5 2 3 4 5
```



```
# Extract the 3rd element in vec4
vec4[3]
```

```
## [1] 4
```

```
# Extract the 3rd through 5th elements in vec4
vec4[3:5]
```

```
## [1] 4 5 2
```

```
# Extract the 1st, 3rd, and 7th elements in vec4
vec4[c(1,3,7)]
```

```
## [1] 2 4 4
```

```
# Remove the 2nd element from vec4
vec4[-2]
```

```
## [1] 2 4 5 2 3 4 5
```

```
# Remove the 2nd, 4th, and 5th elements from vec4
vec4[-c(2,4,5)]
```

```
## [1] 2 4 3 4 5
```

We can also use the following functions to either a) get a logical vector indicating which values in the vector meet some criterion or b) get indices of values in a vector that meet some criterion.

```
# Logical vector of the same length as vec4
# TRUE wherever elements equal 2; FALSE elsewhere
vec4==2
```

```
## [1] TRUE FALSE FALSE FALSE TRUE FALSE FALSE FALSE
```

```
# Get indices of all values in vec4 that are equal to 2
which(vec4==2)
```

```
## [1] 1 5
```

```
# Get index of the maximum value in vec4
# If the maximum occurs more than once, this returns the first location by default
which.max(vec4)
```

```
## [1] 4
```

```
# Return logical vector indicating which elements of vec4 are either equal to 2 or 4
vec4 %in% c(2,4)
```

```
## [1] TRUE FALSE TRUE FALSE TRUE FALSE TRUE FALSE
```

```
# Another way to do the same. Note that | means "or" and & means "and"
vec4==2 | vec4==4
```

```
## [1] TRUE FALSE TRUE FALSE TRUE FALSE TRUE FALSE
```

By enclosing the output from the above functions in square brackets, we can extract elements meeting particular criteria from a vector. For example:

```
# Extract elements of vec4 that are equal to 2
vec4[vec4==2]
```

```
## [1] 2 2
```

```
# Extract elements of vec4 that are equal to 2 or 4  
vec4[vec4==2|vec4==4]
```

```
## [1] 2 4 2 4
```

```
#or:
```

```
vec4[vec4 %in% c(2,4)]
```

```
## [1] 2 4 2 4
```

Using indices to extract elements in a matrix